

Hash Tables

CS 2860: Algorithms and Complexity

Magnus Wahlström and Gregory Gutin

February 10, 2018

Recap: Hash tables

- ▶ Hash tables: “Almost perfect” data structure for unordered Sets/Maps
- ▶ Map: Stores mappings (associations) ($\text{key} \mapsto \text{value}$)
- ▶ **Constant** ($O(1)$) average time for basic operations (get, put, remove, membership test)
- ▶ Does not support any ordering operations (no fast min, no fast successor, iteration in “unknown order” only)

The plan

- ▶ Saw:
 - ▶ Implementation
 - ▶ Hash buckets, collision strategies
- ▶ Today:
 - ▶ Hash functions
 - ▶ The hashCode/equals contract
 - ▶ Other hash applications

Hash Functions

Hash functions

- ▶ **Cryptographic** hash functions – convert file (sequence/block of bytes) to “fingerprint” bitstring
 - ▶ Different inputs should **practically always** give different hashes
 - ▶ Ideally **infeasible** to find even **single pair** with same hash value
 - ▶ Ideally **uninvertible**: Cannot guess input from hash value; data fully scrambled, differences exaggerated
- ▶ Applications:
 - ▶ Digital signatures: use short signature from trusted source to verify nature of large files from untrusted sources
 - ▶ Detect whether file has been changed/corrupted
 - ▶ File identifiers – git, svn, backup systems

Hash functions

- ▶ Our hashes are not cryptographic
- ▶ Compared to cryptographic hashes, what we need is weaker:
 - ▶ Our hashes are 32 bits, not (say) 256 bits
 - ▶ Collisions are annoying, but no disasters
 - ▶ Computation should be fast
- ▶ Another difference:
 - ▶ Cryptographic hash depends on **bit-level raw data**
 - ▶ Our hashes depend on **logical structure** of data
 - ▶ (See **hashCode>equals contract**, later)

Hash functions

- ▶ Our hashes are not cryptographic
- ▶ Compared to cryptographic hashes, what we need is weaker:
 - ▶ Our hashes are 32 bits, not (say) 256 bits
 - ▶ Collisions are annoying, but no disasters
 - ▶ Computation should be fast
- ▶ Another difference:
 - ▶ Cryptographic hash depends on **bit-level raw data**
 - ▶ Our hashes depend on **logical structure** of data
 - ▶ (See **hashCode>equals contract**, later)

Hash functions

- ▶ Our hashes are not cryptographic
- ▶ Compared to cryptographic hashes, what we need is weaker:
 - ▶ Our hashes are 32 bits, not (say) 256 bits
 - ▶ Collisions are annoying, but no disasters
 - ▶ Computation should be fast
- ▶ Another difference:
 - ▶ Cryptographic hash depends on **bit-level raw data**
 - ▶ Our hashes depend on **logical structure** of data
 - ▶ (See **hashCode>equals contract**, later)

Hash functions

- ▶ Our hashes are not cryptographic
- ▶ Compared to cryptographic hashes, what we need is weaker:
 - ▶ Our hashes are 32 bits, not (say) 256 bits
 - ▶ Collisions are annoying, but no disasters
 - ▶ Computation should be fast
- ▶ Another difference:
 - ▶ Cryptographic hash depends on **bit-level raw data**
 - ▶ Our hashes depend on **logical structure** of data
 - ▶ (See **hashCode>equals contract**, later)

★ Large numbers

- ▶ The number of 32-bit hashes is 2^{32}
- ▶ For most types of data, collisions are unavoidable:
 - ▶ The number of 10-character strings made from letters a–z is

$$26^{10} \approx 2^{47}$$

- ▶ There are at least $2^{15} > 30,000$ different 10-letter words that all map to the same hash code (no matter what hash function)
- ▶ The number of files of 1,024 bytes is

$$256^{1024} = 2^{8192}$$

- ▶ The number of different 1K-files with the same hash code is unimaginably large, no matter what hash function you use
- ▶ Most of these strings/files/... will never be created
 1. If we have an **adversary**, they may **maliciously craft** colliding strings (e.g., as a part of a DoS-attack against a web server)
 2. If we're not worried about that scenario, we'll be fine

★ Large numbers

- ▶ The number of 32-bit hashes is 2^{32}
- ▶ For most types of data, collisions are unavoidable:
 - ▶ The number of 10-character strings made from letters a–z is

$$26^{10} \approx 2^{47}$$

- ▶ There are at least $2^{15} > 30,000$ different 10-letter words that all map to the same hash code (no matter what hash function)
- ▶ The number of files of 1,024 bytes is

$$256^{1024} = 2^{8192}$$

- ▶ The number of different 1K-files with the same hash code is unimaginably large, no matter what hash function you use
- ▶ Most of these strings/files/... will never be created
 1. If we have an **adversary**, they may **maliciously craft** colliding strings (e.g., as a part of a DoS-attack against a web server)
 2. If we're not worried about that scenario, we'll be fine

★ Large numbers

- ▶ The number of 32-bit hashes is 2^{32}
- ▶ For most types of data, collisions are unavoidable:
 - ▶ The number of 10-character strings made from letters a–z is

$$26^{10} \approx 2^{47}$$

- ▶ There are at least $2^{15} > 30,000$ different 10-letter words that all map to the same hash code (no matter what hash function)
- ▶ The number of files of 1,024 bytes is

$$256^{1024} = 2^{8192}$$

- ▶ The number of different 1K-files with the same hash code is unimaginably large, no matter what hash function you use
- ▶ Most of these strings/files/... will never be created
 1. If we have an **adversary**, they may **maliciously craft** colliding strings (e.g., as a part of a DoS-attack against a web server)
 2. If we're not worried about that scenario, we'll be fine

★ Large numbers

- ▶ The number of 32-bit hashes is 2^{32}
- ▶ For most types of data, collisions are unavoidable:
 - ▶ The number of 10-character strings made from letters a–z is

$$26^{10} \approx 2^{47}$$

- ▶ There are at least $2^{15} > 30,000$ different 10-letter words that all map to the same hash code (no matter what hash function)
- ▶ The number of files of 1,024 bytes is

$$256^{1024} = 2^{8192}$$

- ▶ The number of different 1K-files with the same hash code is unimaginably large, no matter what hash function you use
- ▶ Most of these strings/files/... will never be created
 1. If we have an **adversary**, they may **maliciously craft** colliding strings (e.g., as a part of a DoS-attack against a web server)
 2. If we're not worried about that scenario, we'll be fine

★ Large numbers

- ▶ The number of 32-bit hashes is 2^{32}
- ▶ For most types of data, collisions are unavoidable:
 - ▶ The number of 10-character strings made from letters a–z is

$$26^{10} \approx 2^{47}$$

- ▶ There are at least $2^{15} > 30,000$ different 10-letter words that all map to the same hash code (no matter what hash function)
- ▶ The number of files of 1,024 bytes is

$$256^{1024} = 2^{8192}$$

- ▶ The number of different 1K-files with the same hash code is unimaginably large, no matter what hash function you use
- ▶ Most of these strings/files/... will never be created
 1. If we have an **adversary**, they may **maliciously craft** colliding strings (e.g., as a part of a DoS-attack against a web server)
 2. If we're not worried about that scenario, we'll be fine

★ Behaviour of random keys

- ▶ Assume a table of capacity m , containing n random items
- ▶ At what value of n does each of the following occur?
 1. The first collision:
 2. Table half-full:
 3. Table 90% full:
 4. Table entirely full:

★ Behaviour of random keys

- ▶ Assume a table of capacity m , containing n random items
- ▶ At what value of n does each of the following occur?
 1. The first collision: At $n = \Omega(\sqrt{m})$ (birthday paradox)
 2. Table half-full:
 3. Table 90% full:
 4. Table entirely full:

★ Behaviour of random keys

- ▶ Assume a table of capacity m , containing n random items
- ▶ At what value of n does each of the following occur?
 1. The first collision: At $n = \Omega(\sqrt{m})$ (birthday paradox)
 2. Table half-full: At 70% use ($n = 0.7m$)
 3. Table 90% full:
 4. Table entirely full:

★ Behaviour of random keys

- ▶ Assume a table of capacity m , containing n random items
- ▶ At what value of n does each of the following occur?
 1. The first collision: At $n = \Omega(\sqrt{m})$ (birthday paradox)
 2. Table half-full: At 70% use ($n = 0.7m$)
 3. Table 90% full: At $\approx 200\%$ use
 4. Table entirely full:

★ Behaviour of random keys

- ▶ Assume a table of capacity m , containing n random items
- ▶ At what value of n does each of the following occur?
 1. The first collision: At $n = \Omega(\sqrt{m})$ (birthday paradox)
 2. Table half-full: At 70% use ($n = 0.7m$)
 3. Table 90% full: At $\approx 200\%$ use
 4. Table entirely full: At $n = \Omega(m \log m)$ (coupon collection)

★ Behaviour of random keys

- ▶ Assume a table of capacity m , containing n random items
- ▶ At what value of n does each of the following occur?
 1. The first collision: At $n = \Omega(\sqrt{m})$ (birthday paradox)
 2. Table half-full: At 70% use ($n = 0.7m$)
 3. Table 90% full: At $\approx 200\%$ use
 4. Table entirely full: At $n = \Omega(m \log m)$ (coupon collection)
- ▶ When $n = m$, the fraction of free slots is approximately

$$\left(1 - \frac{1}{m}\right)^m = \left(1 - \frac{1}{m}\right)^m \approx 1/e \approx 37\%$$

★ Behaviour of random keys

- ▶ Assume a table of capacity m , containing n random items
- ▶ At what value of n does each of the following occur?
 1. The first collision: At $n = \Omega(\sqrt{m})$ (birthday paradox)
 2. Table half-full: At 70% use ($n = 0.7m$)
 3. Table 90% full: At $\approx 200\%$ use
 4. Table entirely full: At $n = \Omega(m \log m)$ (coupon collection)
- ▶ When $n = m$, the fraction of free slots is approximately

$$(1 - \frac{1}{m})^n = (1 - \frac{1}{m})^m \approx 1/e \approx 37\%$$

- ▶ The “longest run” when $n = m$ is probably $O(\log n)$, but most slots are short ($0/1/2$ elements)

★ Behaviour of random keys

- ▶ Assume a table of capacity m , containing n random items
- ▶ At what value of n does each of the following occur?
 1. The first collision: At $n = \Omega(\sqrt{m})$ (birthday paradox)
 2. Table half-full: At 70% use ($n = 0.7m$)
 3. Table 90% full: At $\approx 200\%$ use
 4. Table entirely full: At $n = \Omega(m \log m)$ (coupon collection)
- ▶ When $n = m$, the fraction of free slots is approximately

$$(1 - \frac{1}{m})^n = (1 - \frac{1}{m})^m \approx 1/e \approx 37\%$$

- ▶ The “longest run” when $n = m$ is probably $O(\log n)$, but most slots are short ($0/1/2$ elements)
- ▶ Common strategy: Move to larger table at $n = 0.7m$

Desired properties of “our” hash functions

- ▶ Easy to compute (e.g., $O(n)$ time if object has size n)
- ▶ Reasonably irregular (few collisions)
- ▶ Few collisions even in the face of realistic, non-random use patterns
- ▶ Consistent with equals

Desired properties of “our” hash functions

- ▶ Easy to compute (e.g., $O(n)$ time if object has size n)
- ▶ Reasonably irregular (few collisions)
- ▶ Few collisions even in the face of realistic, non-random use patterns
- ▶ Consistent with equals

Desired properties of “our” hash functions

- ▶ Easy to compute (e.g., $O(n)$ time if object has size n)
- ▶ Reasonably irregular (few collisions)
- ▶ Few collisions even in the face of **realistic, non-random** use patterns
- ▶ Consistent with **equals**

Desired properties of “our” hash functions

- ▶ Easy to compute (e.g., $O(n)$ time if object has size n)
- ▶ Reasonably irregular (few collisions)
- ▶ Few collisions even in the face of **realistic, non-random** use patterns
- ▶ Consistent with **equals**

The hashCode/equals contract

hashCode and equals

- ▶ Fact: All objects in Java have the following two methods:
 - ▶ `object.hashCode()`: return hash value
 - ▶ `object1.equals(object2)`: equivalence test
- ▶ There's an important **contract** between these two:
 1. If you change one of these methods (e.g., `equals`), you also need to change the other!
 2. If you do, your methods must be **consistent** with each other!
- ▶ Why?

hashCode and equals

- ▶ Fact: All objects in Java have the following two methods:
 - ▶ `object.hashCode()`: return hash value
 - ▶ `object1.equals(object2)`: equivalence test
- ▶ There's an important **contract** between these two:
 1. If you change one of these methods (e.g., equals), you also need to change the other!
 2. If you do, your methods must be **consistent** with each other!
- ▶ Why?

hashCode and equals

- ▶ Fact: All objects in Java have the following two methods:
 - ▶ `object.hashCode()`: return hash value
 - ▶ `object1.equals(object2)`: equivalence test
- ▶ There's an important **contract** between these two:
 1. If you change one of these methods (e.g., equals), you also need to change the other!
 2. If you do, your methods must be **consistent** with each other!
- ▶ Why?
- ▶ Consistency:
 - ▶ If `x.equals(y)` then `x.hashCode()==y.hashCode()`!
 - ▶ Converse: If `x.hashCode()!=y.hashCode()`, then `!x.equals(y)`.

Testing object equivalence

- ▶ Recall: Want to test (for Strings, Sets, Arrays...) whether two **distinct** objects have **equivalent** contents
- ▶ Example: Two different strings both contain "rabbit"
- ▶ Example: Two different **Sets** contain the same items (but in different orders in the tree/in the hash buckets!)
- ▶ So to test whether two **Sets** are equal, we cannot use "bit-level identity" (as a cryptographic hash code would give us). We have to do:

Testing object equivalence

- ▶ Recall: Want to test (for Strings, Sets, Arrays...) whether two **distinct** objects have **equivalent** contents
- ▶ Example: Two different strings both contain "rabbit"
- ▶ Example: Two different **Sets** contain the same items (but in different orders in the tree/in the hash buckets!)
- ▶ So to test whether two **Sets** are equal, we cannot use "bit-level identity" (as a cryptographic hash code would give us). We have to do:

Testing object equivalence

- ▶ Recall: Want to test (for Strings, Sets, Arrays...) whether two **distinct** objects have **equivalent** contents
- ▶ Example: Two different strings both contain "rabbit"
- ▶ Example: Two different **Sets** contain the same items (but in different orders in the tree/in the hash buckets!)
- ▶ So to test whether two **Sets** are equal, we cannot use "bit-level identity" (as a cryptographic hash code would give us). We have to do:
 1. If **set1.size** **!=** **set2.size**, they are different
 2. Otherwise, for every object **x** in set1:
 - 2.1 If set2 does not contain **x** (or an object equivalent to **x**!) the sets are different¹
 3. If the loop terminates, the sets are equivalent

¹The test "contains an object equivalent to **x**" is just **set2.contains(x)**

Testing object equivalence

- ▶ Recall: Want to test (for Strings, Sets, Arrays...) whether two **distinct** objects have **equivalent** contents
- ▶ Example: Two different strings both contain "rabbit"
- ▶ Example: Two different **Sets** contain the same items (but in different orders in the tree/in the hash buckets!)
- ▶ So to test whether two **Sets** are equal, we cannot use "bit-level identity" (as a cryptographic hash code would give us). We have to do:
 1. If **set1.size** **!=** **set2.size**, they are different
 2. Otherwise, for every object **x** in set1:
 - 2.1 If set2 does not contain **x** (or an object equivalent to **x**!) the sets are different¹
 3. If the loop terminates, the sets are equivalent
- ▶ So since a Set must implement its own **equals** code, it must also implement a new **hashCode** function

¹The test "contains an object equivalent to **x**" is just **set2.contains(x)**

Default operations

- ▶ The default rules (for Object) are:
 - ▶ object.hashCode() returns the **physical memory address**
 - ▶ object1.equals(object2) tests **object1 == object2?**
- ▶ Consistent?

Default operations

- ▶ The default rules (for Object) are:
 - ▶ `object.hashCode()` returns the `physical memory address`
 - ▶ `object1.equals(object2)` tests `object1 == object2?`
- ▶ Consistent?

Default operations

- ▶ The default rules (for Object) are:
 - ▶ `object.hashCode()` returns the `physical memory address`
 - ▶ `object1.equals(object2)` tests `object1 == object2?`
- ▶ Consistent? Yes.

Default operations

- ▶ The default rules (for Object) are:
 - ▶ `object.hashCode()` returns the **physical memory address**
 - ▶ `object1.equals(object2)` tests **`object1 == object2`**?
- ▶ Consistent? Yes.
- ▶ But it's also a very “boring” equality test.
- ▶ If you subclass directly from Object, this is your default functionality

A bad hash function for Strings

- ▶ Recall:
 1. **String** objects test equality by character-by-character comparisons
 2. Every character in a String maps to a **unicode code page** (a number $0-2^{32}$, usually $0-255$)
- ▶ What about the following hash code for strings?
 1. Initialise `sum=0`
 2. For every character `c` in the string:
 - 2.1 Add the **numerical value** of `c` to `sum`
 3. Return resulting `sum`
- ▶ Is it (1) consistent? (2) good?

A bad hash function for Strings

- ▶ Recall:
 1. **String** objects test equality by character-by-character comparisons
 2. Every character in a String maps to a **unicode code page** (a number $0-2^{32}$, usually $0-255$)
- ▶ What about the following hash code for strings?
 1. Initialise $\text{sum}=0$
 2. For every character **c** in the string:
 - 2.1 Add the **numerical value** of **c** to sum
 3. Return resulting sum
- ▶ Is it (1) consistent? (2) good?

A bad hash function for Strings

- ▶ Recall:
 1. **String** objects test equality by character-by-character comparisons
 2. Every character in a String maps to a **unicode code page** (a number $0-2^{32}$, usually $0-255$)
- ▶ What about the following hash code for strings?
 1. Initialise $\text{sum}=0$
 2. For every character **c** in the string:
 - 2.1 Add the **numerical value** of **c** to sum
 3. Return resulting sum
- ▶ Is it (1) consistent? (2) good?

A bad hash function for Strings

- ▶ Recall:
 1. **String** objects test equality by character-by-character comparisons
 2. Every character in a String maps to a **unicode code page** (a number $0-2^{32}$, usually $0-255$)
- ▶ What about the following hash code for strings?
 1. Initialise $\text{sum}=0$
 2. For every character **c** in the string:
 - 2.1 Add the **numerical value** of **c** to sum
 3. Return resulting sum
- ▶ Is it (1) consistent? (2) good?
 1. Yes: Character-by-character equality tests on bit-level.

A bad hash function for Strings

- ▶ Recall:
 1. **String** objects test equality by character-by-character comparisons
 2. Every character in a String maps to a **unicode code page** (a number $0-2^{32}$, usually $0-255$)
- ▶ What about the following hash code for strings?
 1. Initialise $\text{sum}=0$
 2. For every character **c** in the string:
 - 2.1 Add the **numerical value** of **c** to sum
 3. Return resulting sum
- ▶ Is it (1) consistent? (2) good?
 1. Yes: Character-by-character equality tests on bit-level.
 2. High quality?
 - ▶ $\text{hash}(\text{god}) == \text{hash}(\text{dog}), \text{hash}(\text{cat}) == \text{hash}(\text{act})$
 - ▶ $\text{hash}(\text{reset}) == \text{hash}(\text{trees}) == \text{hash}(\text{steer}) == \text{hash}(\text{terse})$

A bad hash function for Strings

- ▶ Recall:
 1. **String** objects test equality by character-by-character comparisons
 2. Every character in a String maps to a **unicode code page** (a number $0-2^{32}$, usually $0-255$)
- ▶ What about the following hash code for strings?
 1. Initialise $sum=0$
 2. For every character **c** in the string:
 - 2.1 Add the **numerical value** of **c** to sum
 3. Return resulting sum
- ▶ Is it (1) consistent? (2) good?
 1. Yes: Character-by-character equality tests on bit-level.
 2. High quality?
 - ▶ $hash(god) == hash(dog), hash(cat) == hash(act)$
 - ▶ $hash(reset) == hash(trees) == hash(steer) == hash(terse)$
- ▶ Hash code is low-quality on **common use patterns**, because English text contains **many anagrams**
- ▶ Also, it “throws away” the dependency on order

Immutable objects

- ▶ A common programming trick is to make objects **immutable** – unchangeable
- ▶ Example String objects:
 - ▶ Cannot **change** one character inside a String
 - ▶ Can return a **different String** with one character changed
- ▶ Consequence for hashes:
 - ▶ Once we have computed a hashCode value, we may **remember** the hash value in a member variable and never compute it again
- ▶ (Caching hash values is also possible for **mutable** objects, if you're careful to **forget** the hash every time it changes)
- ▶ Note: Computing hashCode should take at least $\Omega(n)$ time (but hopefully not more)

Immutable objects

- ▶ A common programming trick is to make objects **immutable** – unchangeable
- ▶ Example String objects:
 - ▶ Cannot **change** one character inside a String
 - ▶ Can return a **different String** with one character changed
- ▶ Consequence for hashes:
 - ▶ Once we have computed a hashCode value, we may **remember** the hash value in a member variable and never compute it again
- ▶ (Caching hash values is also possible for **mutable** objects, if you're careful to **forget** the hash every time it changes)
- ▶ Note: Computing hashCode should take at least $\Omega(n)$ time (but hopefully not more)

Immutable objects

- ▶ A common programming trick is to make objects **immutable** – unchangeable
- ▶ Example String objects:
 - ▶ Cannot **change** one character inside a String
 - ▶ Can return a **different String** with one character changed
- ▶ Consequence for hashes:
 - ▶ Once we have computed a hashCode value, we may **remember** the hash value in a member variable and never compute it again
- ▶ (Caching hash values is also possible for **mutable** objects, if you're careful to **forget** the hash every time it changes)
- ▶ Note: Computing hashCode should take at least $\Omega(n)$ time (but hopefully not more)

Immutable objects

- ▶ A common programming trick is to make objects **immutable** – unchangeable
- ▶ Example String objects:
 - ▶ Cannot **change** one character inside a String
 - ▶ Can return a **different String** with one character changed
- ▶ Consequence for hashes:
 - ▶ Once we have computed a hashCode value, we may **remember** the hash value in a member variable and never compute it again
- ▶ (Caching hash values is also possible for **mutable** objects, if you're careful to **forget** the hash every time it changes)
- ▶ Note: Computing hashCode should take at least $\Omega(n)$ time (but hopefully not more)

Mutable/immutable

- ▶ Suggestions?

Mutable/immutable

- ▶ Suggestions?
- ▶ Necessarily mutable: Array, Set, Map...

Mutable/immutable

- ▶ Suggestions?
- ▶ Necessarily mutable: Array, Set, Map. . .
- ▶ Easily immutable: Single-linked lists (with some programming discipline)
- ▶ Same principle for binary trees **without parent links**

Mutable key types

Warning: Never modify an object after it has been used as the key of a hash table!

- ▶ Many objects (String, Integer, ...) are **immutable**: They cannot be modified after creation
 - ▶ You cannot change the value of 3 to 5, only “rebind” a single variable from 3 to 5; similarly with strings
- ▶ However, some modifiable data types (e.g., arrays/lists) are still usable as hash table key types
- ▶ Modifying these after insertion is a bug!

Mutable key types

Warning: Never modify an object after it has been used as the key of a hash table!

- ▶ Many objects (String, Integer, ...) are **immutable**: They cannot be modified after creation
 - ▶ You cannot change the value of 3 to 5, only “rebind” a single variable from 3 to 5; similarly with strings
- ▶ However, some modifiable data types (e.g., arrays/lists) are still usable as hash table key types
- ▶ Modifying these after insertion is a bug!

Mutable key types

Warning: Never modify an object after it has been used as the key of a hash table!

- ▶ Many objects (String, Integer, ...) are **immutable**: They cannot be modified after creation
 - ▶ You cannot change the value of 3 to 5, only “rebind” a single variable from 3 to 5; similarly with strings
- ▶ However, some modifiable data types (e.g., arrays/lists) are still usable as hash table key types
- ▶ Modifying these after insertion is a bug!

Mutable key types

Warning: Never modify an object after it has been used as the key of a hash table!

- ▶ Many objects (String, Integer, ...) are **immutable**: They cannot be modified after creation
 - ▶ You cannot change the value of 3 to 5, only “rebind” a single variable from 3 to 5; similarly with strings
- ▶ However, some modifiable data types (e.g., arrays/lists) are still usable as hash table key types
- ▶ Modifying these after insertion is a bug!

Immutable keys warning explained

- ▶ Lookup procedure: Consider call `lookup(item)`
 1. Let `hash1 = item.hashCode()`
 2. Compute hash bucket `bucket = (hash1 % m)`
(table capacity `m`)
 3. For each element `item2` stored in slot `bucket`:
 - 3.1 If `item2.hashCode() != hash1`: Continue to next object
 - 3.2 If `item2.hashCode() == hash1`, test `item2.equals(item)`
 - 3.3 If positive, return YES (found a copy of item in the hash table)
- ▶ What happens if the “old” item `item2` was modified after insertion?
 - ▶ The new and the old versions can have **different hash codes**:
The item would be in the wrong place!
 - ▶ Even a test `hashtable.contains(item2)` would fail after modification!

Immutable keys warning explained

- ▶ Lookup procedure: Consider call `lookup(item)`
 1. Let `hash1 = item.hashCode()`
 2. Compute hash bucket `bucket = (hash1 % m)`
(table capacity `m`)
 3. For each element `item2` stored in slot `bucket`:
 - 3.1 If `item2.hashCode() != hash1`: Continue to next object
 - 3.2 If `item2.hashCode() == hash1`, test `item2.equals(item)`
 - 3.3 If positive, return YES (found a copy of item in the hash table)
- ▶ What happens if the “old” item `item2` was modified after insertion?
 - ▶ The new and the old versions can have **different hash codes**:
The item would be in the wrong place!
 - ▶ Even a test `hashtable.contains(item2)` would fail after modification!

Hash functions in Java

- ▶ Integer: $\text{hash}(i)=i$
- ▶ Single character: $\text{hash}(c)=\text{Unicode code point of } c$
- ▶ String S , length n :
 - ▶ $s[0] \cdot 31^{n-1} + s[1] \cdot 31^{n-2} + \dots + s[n-1] \cdot 1$
- ▶ Array/ordered collection A :
 - ▶ $\text{hash}(A[0]) \cdot 31^{n-1} + \text{hash}(A[1]) \cdot 31^{n-2} + \dots$
- ▶ Unordered collection (set) S :
 - ▶ Sum of $\text{hash}(x)$ over all x in S

Hash functions in Java

- ▶ Integer: $\text{hash}(i)=i$
- ▶ Single character: $\text{hash}(c)=\text{Unicode code point of } c$
- ▶ String S , length n :
 - ▶ $s[0] \cdot 31^{n-1} + s[1] \cdot 31^{n-2} + \dots + s[n-1] \cdot 1$
- ▶ Array/ordered collection A :
 - ▶ $\text{hash}(A[0]) \cdot 31^{n-1} + \text{hash}(A[1]) \cdot 31^{n-2} + \dots$
- ▶ Unordered collection (set) S :
 - ▶ Sum of $\text{hash}(x)$ over all x in S

Hash functions in Java

- ▶ Integer: $\text{hash}(i)=i$
- ▶ Single character: $\text{hash}(c)=\text{Unicode code point of } c$
- ▶ String S , length n :
 - ▶ $s[0] \cdot 31^{n-1} + s[1] \cdot 31^{n-2} + \dots + s[n-1] \cdot 1$
- ▶ Array/ordered collection A :
 - ▶ $\text{hash}(A[0]) \cdot 31^{n-1} + \text{hash}(A[1]) \cdot 31^{n-2} + \dots$
- ▶ Unordered collection (set) S :
 - ▶ Sum of $\text{hash}(x)$ over all x in S

Hash functions in Java

- ▶ Integer: $\text{hash}(i)=i$
- ▶ Single character: $\text{hash}(c)=\text{Unicode code point of } c$
- ▶ String S , length n :
 - ▶ $s[0] \cdot 31^{n-1} + s[1] \cdot 31^{n-2} + \dots + s[n-1] \cdot 1$
- ▶ Array/ordered collection A :
 - ▶ $\text{hash}(A[0]) \cdot 31^{n-1} + \text{hash}(A[1]) \cdot 31^{n-2} + \dots$
- ▶ Unordered collection (set) S :
 - ▶ Sum of $\text{hash}(x)$ over all x in S

Other hash applications (v. briefly)

Applications

1. To **speed up** equivalence testing (if hashCode has been cached)
2. More clever **string comparison** (Rabin-Karp algorithm)
 - 2.1 Compute hash(pattern) length m
 - 2.2 Compute hash(string[0 ... $m-1$]); compare hash values
 - 2.3 **Update** to hash(string[1 ... m]); compare values
 - 2.4 **Update** to hash(string[2 ... m]); compare values ...
3. **Bloom filters**
 - 3.1 When working with **very large** objects, instead of a full hash table, implement only the **Boolean array** solution
 - 3.2 Can (probabilistically!) answer question “have I seen this object before?”
 - 3.3 Space $O(n)$ for n objects of size S , much better than $n \cdot S$
 - 3.4 Example: “Does computer ... in my distributed network have a copy of this file?”

Applications

1. To **speed up** equivalence testing (if hashCode has been cached)
2. More clever **string comparison** (Rabin-Karp algorithm)
 - 2.1 Compute hash(pattern) length m
 - 2.2 Compute hash(string[0 ... $m-1$]); compare hash values
 - 2.3 **Update** to hash(string[1 ... m]); compare values
 - 2.4 **Update** to hash(string[2 ... m]); compare values ...
3. **Bloom filters**
 - 3.1 When working with **very large** objects, instead of a full hash table, implement only the **Boolean array** solution
 - 3.2 Can (probabilistically!) answer question “have I seen this object before?”
 - 3.3 Space $O(n)$ for n objects of size S , much better than $n \cdot S$
 - 3.4 Example: “Does computer ... in my distributed network have a copy of this file?”

Applications

1. To **speed up** equivalence testing (if hashCode has been cached)
2. More clever **string comparison** (Rabin-Karp algorithm)
 - 2.1 Compute hash(pattern) length m
 - 2.2 Compute hash(string[0 ... $m-1$]); compare hash values
 - 2.3 **Update** to hash(string[1 ... m]); compare values
 - 2.4 **Update** to hash(string[2 ... m]); compare values ...
3. **Bloom filters**
 - 3.1 When working with **very large** objects, instead of a full hash table, implement only the **Boolean array** solution
 - 3.2 Can (probabilistically!) answer question “have I seen this object before?”
 - 3.3 Space $O(n)$ for n objects of size S , much better than $n \cdot S$
 - 3.4 Example: “Does computer ... in my distributed network have a copy of this file?”

Applications

1. To **speed up** equivalence testing (if hashCode has been cached)
2. More clever **string comparison** (Rabin-Karp algorithm)
 - 2.1 Compute hash(pattern) length m
 - 2.2 Compute hash(string[0 ... $m-1$]); compare hash values
 - 2.3 **Update** to hash(string[1 ... m]); compare values
 - 2.4 **Update** to hash(string[2 ... m]); compare values ...
3. **Bloom filters**
 - 3.1 When working with **very large** objects, instead of a full hash table, implement only the **Boolean array** solution
 - 3.2 Can (probabilistically!) answer question “have I seen this object before?”
 - 3.3 Space $O(n)$ for n objects of size S , much better than $n \cdot S$
 - 3.4 Example: “Does computer ... in my distributed network have a copy of this file?”

Applications

1. To **speed up** equivalence testing (if hashCode has been cached)
2. More clever **string comparison** (Rabin-Karp algorithm)
 - 2.1 Compute hash(pattern) length m
 - 2.2 Compute hash(string[0 ... $m-1$]); compare hash values
 - 2.3 **Update** to hash(string[1 ... m]); compare values
 - 2.4 **Update** to hash(string[2 ... m]); compare values ...
3. **Bloom filters**
 - 3.1 When working with **very large** objects, instead of a full hash table, implement only the **Boolean array** solution
 - 3.2 Can (probabilistically!) answer question “have I seen this object before?”
 - 3.3 Space $O(n)$ for n objects of size S , much better than $n \cdot S$
 - 3.4 Example: “Does computer ... in my distributed network have a copy of this file?”

Summary

We saw:

- ▶ Properties of hash functions and random numbers
- ▶ The hashCode/equals contract
- ▶ A few samples of “other algorithmic applications” of hash functions (may return to these)